

# Contents

<b>1 Chapter 4: Sparse Reach + Push -- Introduce HER Where It Matters</b>	<b>2</b>
1.1 WHY: The Sparse Reward Problem . . . . .	2
1.1.1 What "Sparse" Means in Gymnasium-Robotics . . . . .	2
1.1.2 Why Standard RL Fails on Sparse Goals . . . . .	2
1.1.3 The Key Insight: Failed Trajectories Contain Information . . . . .	2
1.2 HOW: Hindsight Experience Replay (HER) . . . . .	3
1.2.1 The Relabeling Strategy . . . . .	3
1.2.2 The <code>n_sampled_goal</code> Parameter . . . . .	3
1.2.3 Why HER Requires Off-Policy Learning . . . . .	3
1.3 BUILD IT: HER Relabeling in Code . . . . .	3
1.3.1 Goal Sampling Strategies . . . . .	4
1.3.2 Relabeling a Transition . . . . .	4
1.3.3 The Data Amplification Effect . . . . .	4
1.3.4 Verify the Lab . . . . .	4
1.3.5 Exercises: Modify and Observe . . . . .	4
1.4 WHAT: Experiments and Expected Results (Run It) . . . . .	5
1.4.1 Running the Experiments . . . . .	5
1.4.2 Experiment 1: Sparse Reach -- HER vs No-HER . . . . .	6
1.4.3 Experiment 2: Sparse Push -- Where HER Matters . . . . .	7
1.4.4 Sweep Results: 120 Runs, 24 Configurations . . . . .	9
1.4.5 Mathematical Analysis: Why Gamma Dominates . . . . .	11
1.4.6 Visual Comparison: Best vs Worst Policy . . . . .	12
1.4.7 Experiment 3: Full Factorial Ablation (The Sweep) . . . . .	16
1.4.8 CLI Parameter Reference . . . . .	16
1.5 Monitoring and Debugging . . . . .	18
1.5.1 Artifact Locations . . . . .	18
1.5.2 Checking Training Progress . . . . .	18
1.5.3 TensorBoard Metrics to Watch . . . . .	18
1.5.4 Common Issues . . . . .	18
1.6 Interpreting Results . . . . .	19
1.6.1 What "Clear Separation" Means . . . . .	19
1.6.2 Interpreting Our Results . . . . .	19
1.6.3 Statistical Validity . . . . .	20
1.6.4 What to Do Next . . . . .	20
1.7 The Reality of RL Research: Debugging and Experimentation . . . . .	20
1.7.1 Why Your First Attempt Will Probably Fail . . . . .	20
1.7.2 Our Diagnostic Journey . . . . .	20
1.7.3 How to Debug Your Own RL Experiments . . . . .	22
1.7.4 The Metrics We Used for Diagnosis . . . . .	23
1.7.5 Lessons Learned . . . . .	23
1.7.6 Time Investment Reality Check . . . . .	24
1.7.7 The Plot Twist: Our 98% Was a Lucky Seed . . . . .	24
1.7.8 Checkpoint Availability . . . . .	24
1.8 Summary . . . . .	25
1.8.1 Key Findings from Our Experiments . . . . .	25
1.8.2 Sweep-Validated Recommendations . . . . .	25
1.8.3 Files Generated . . . . .	26
1.9 References . . . . .	26

# 1 Chapter 4: Sparse Reach + Push -- Introduce HER Where It Matters

**Week 4 Goal:** Demonstrate that Hindsight Experience Replay (HER) is the difference-maker on sparse goal-conditioned tasks.

---

## 1.1 WHY: The Sparse Reward Problem

In Chapter 3, we trained SAC on FetchReachDense-v4, where every timestep provides a shaped reward signal proportional to the distance to the goal. The agent received continuous feedback: "you're getting warmer" or "you're getting colder." This made learning straightforward.

Now we face the real challenge: **sparse rewards**.

### 1.1.1 What "Sparse" Means in Gymnasium-Robotics

In sparse Fetch environments (FetchReach-v4, FetchPush-v4, FetchPickAndPlace-v4), the reward function is:

$$R(s, a, g) = \begin{cases} 0 & \text{if } \|g_{\text{achieved}} - g_{\text{desired}}\| < \epsilon \\ -1 & \text{otherwise} \end{cases}$$

where  $\epsilon = 0.05$  meters (5 cm) is the success threshold.

**The problem:** The agent receives  $R = -1$  for almost every transition until it accidentally succeeds. With random exploration, this might never happen -- or happen so rarely that the agent cannot learn from it.

### 1.1.2 Why Standard RL Fails on Sparse Goals

Consider training SAC without HER on sparse Reach:

1. **Initial exploration is random.** The gripper moves chaotically.
2. **Most episodes fail.** The gripper rarely lands within 5cm of the goal by chance.
3. **All transitions have reward  $-1$ .** The critic learns: "everything is equally bad."
4. **No gradient signal for improvement.** Without reward variation, the policy has no direction to improve.

This is not a hyperparameter problem -- it is a **structural limitation** of standard RL with sparse rewards.

### 1.1.3 The Key Insight: Failed Trajectories Contain Information

Here is the insight that makes HER work:

A trajectory that fails to reach goal  $g$  is a **successful demonstration** of how to reach wherever it ended up.

If the gripper tried to reach position  $(0.3, 0.2, 0.1)$  but ended at  $(0.5, 0.4, 0.15)$ , we have evidence that the executed actions lead to  $(0.5, 0.4, 0.15)$ . We can **relabel** the trajectory: pretend the goal was  $(0.5, 0.4, 0.15)$  all along, and now we have a successful episode with reward 0.

This is Hindsight Experience Replay.

---

## 1.2 HOW: Hindsight Experience Replay (HER)

### 1.2.1 The Relabeling Strategy

For each transition  $(s_t, a_t, r_t, s_{t+1}, g)$  stored in the replay buffer, HER also stores relabeled versions:

1. **Original transition:** goal =  $g$  (the intended goal)
2. **Relabeled transitions:** goal =  $g'$  (substituted goals)

The substituted goals  $g'$  come from the same trajectory. Common strategies:

Strategy	Description
future	Sample $g'$ from achieved goals at timesteps $t' > t$ in the same episode
final	Use only the final achieved goal $g' = g_{\text{achieved}}(s_T)$
episode	Sample $g'$ from any achieved goal in the episode

**Why future works best:** It creates more relabeled transitions (one for each future timestep), and these goals are "reachable" from the current state -- the agent just demonstrated it can get there.

### 1.2.2 The n\_sampled\_goal Parameter

For each original transition, HER creates `n_sampled_goal` additional relabeled transitions. The default is 4, meaning:

- 1 original transition (goal = intended)
- 4 relabeled transitions (goal = sampled achieved goals)

This 5x expansion of the replay buffer is how HER manufactures dense reward signal from sparse feedback.

### 1.2.3 Why HER Requires Off-Policy Learning

HER only works with **off-policy** algorithms (SAC, TD3, DDPG) because:

1. **Relabeling changes the reward.** The relabeled transition has a different reward than what the agent actually experienced.
2. **Off-policy methods don't care.** They can learn from any transition, regardless of which policy generated it.
3. **On-policy methods (PPO) cannot use relabeled data.** They require transitions from the current policy.

This is why the syllabus builds SAC mastery (Weeks 2-3) before introducing HER (Week 4).

---

## 1.3 BUILD IT: HER Relabeling in Code

This section shows how HER's goal relabeling maps to code. We use pedagogical implementations from `scripts/labs/her_relabeler.py`—these are for understanding, not production.

### 1.3.1 Goal Sampling Strategies

The "future" strategy samples achieved goals from timesteps after the current transition:

```
--8<-- "scripts/labs/her_relabeler.py:goal_sampling"
```

**Key insight:** The future strategy ensures temporal consistency—we only relabel with goals the agent *actually reached* from states similar to the current one.

### 1.3.2 Relabeling a Transition

The core HER operation: substitute a new goal and recompute the reward:

```
--8<-- "scripts/labs/her_relabeler.py:relabel_transition"
```

**Key mapping:**

Concept	Code	Meaning
Original goal	transition.desired_goal	What we were trying to reach
Achieved goal	transition.achieved_goal	Where we actually ended up
Relabeled goal	new_goal	Substitute this as the "desired" goal
New reward	compute_reward_fn(achieved, new_goal)	Did achieved match the new goal?

**Crucial:** The achieved\_goal stays the same—only desired\_goal changes. If achieved == new desired, the transition becomes a "success."

### 1.3.3 The Data Amplification Effect

Processing an episode with HER dramatically increases the success rate in the replay buffer:

```
--8<-- "scripts/labs/her_relabeler.py:her_buffer_insert"
```

**Without HER:** Nearly 0% of transitions have positive reward (sparse signal). **With HER:** Many relabeled transitions are "successes" because achieved == relabeled goal.

### 1.3.4 Verify the Lab

Run the from-scratch implementation's sanity checks:

```
bash docker/dev.sh python scripts/labs/her_relabeler.py --verify
```

Expected output:

- Goal sampling produces correct number of goals
- Relabeling with own achieved goal produces reward=0 (success)
- HER processing increases success rate in synthetic data (0% → ~16%)

This lab is **not** how we train policies—SB3's HER wrapper handles that. The lab shows *what* relabeling does to your data.

### 1.3.5 Exercises: Modify and Observe

#### Exercise: Goal Sampling Strategy Comparison

Run the demo to see how different strategies affect success rates:

```
bash docker/dev.sh python scripts/labs/her_relabeler.py --demo
```

*Observe:* The future, final, and episode strategies produce different success rates. Why does future typically work best?

### Exercise: Relabeling Ratio (k)

In `process_episode_with_her()`, change `k=4` to different values:

# Try: `k=1, k=4, k=8, k=16`

*Question:* How does the success fraction change with `k`? What's the tradeoff between more relabeled data and data quality?

### Exercise: HER Ratio

Change `her_ratio=0.8` to different values:

# Try: `her_ratio=0.0 (no HER), her_ratio=0.5, her_ratio=1.0 (all relabeled)`

*Question:* With `her_ratio=0.0`, you get pure sparse reward learning. With `her_ratio=1.0`, you get maximum relabeling. Why might 0.8 be a good balance?

### Exercise: Understand Why HER Only Gets ~16% Success

The verification shows ~16% success rate on synthetic data. This seems low—why not higher?

*Hint:* The synthetic trajectory uses random actions, so achieved goals are scattered randomly. With a *trained* policy that moves toward goals, future achieved goals would be closer to the current state, making more relabeled transitions successful. HER amplifies competence—it doesn't create it from nothing.

---

## 1.4 WHAT: Experiments and Expected Results (Run It)

### 1.4.1 Running the Experiments

All experiments run through Docker via the `docker/dev.sh` wrapper. This ensures reproducible environments with correct GPU access, MuJoCo rendering, and Python dependencies.

#### 1.4.1.1 Quick Start: Full Pipeline

```
# FetchReach-v4: ~1 hour for all 6 runs (default settings work well)
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py env-all \
    --env FetchReach-v4 --seeds 0,1,2 --total-steps 500000

# FetchPush-v4: ~3 hours for all 6 runs (needs fixed entropy for reliable learning)
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py env-all \
    --env FetchPush-v4 --seeds 0,1,2 --total-steps 2000000 --ent-coef 0.05
```

The `env-all` command runs the complete pipeline:

1. Train SAC without HER (3 seeds)
2. Evaluate each no-HER checkpoint (100 episodes each)
3. Train SAC with HER (3 seeds)
4. Evaluate each HER checkpoint (100 episodes each)
5. Generate comparison report

**Note:** Legacy `reach-all` and `push-all` commands still work for backwards compatibility.

**1.4.1.2 Individual Commands** For more control, run steps separately:

```
# Train a single model
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
    --env FetchPush-v4 --seed 0 --total-steps 500000 # no-HER

bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
    --env FetchPush-v4 --her --seed 0 --total-steps 500000 # with HER

# Train with fixed entropy (recommended for sparse Push)
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
    --env FetchPush-v4 --her --seed 0 --total-steps 2000000 --ent-coef 0.05

# Evaluate a checkpoint
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py eval \
    --ckpt checkpoints/sac_her_FetchPush-v4_seed0.zip

# Compare results across seeds
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py compare \
    --env FetchPush-v4 --seeds 0,1,2
```

**1.4.1.3 Long-Running Jobs with tmux** Training takes 10-15 minutes per run. For the full pipeline (~3 hours for Push), use tmux:

```
# Start a persistent session
tmux new -s week4

# Inside tmux, run the experiments
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py env-all \
    --env FetchPush-v4 --seeds 0,1,2 --ent-coef 0.05 --total-steps 2000000

# Detach: Ctrl-b d
# Reattach later: tmux attach -t week4
```

#### 1.4.1.4 Expected Training Time

Environment	Steps	Time per run	Total (6 runs)
FetchReach-v4	500k	~10 min	~60 min
FetchPush-v4	500k	~14 min	~85 min
FetchPush-v4	1M	~28 min	~170 min

Times measured on DGX with RTX A100, ~600 fps throughput.

---

#### 1.4.2 Experiment 1: Sparse Reach -- HER vs No-HER

**Hypothesis:** SAC without HER will struggle on FetchReach-v4; SAC with HER will succeed more reliably.

```
# Full pipeline: train both, evaluate, compare (3 seeds each)
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py reach-all --seeds 0,1,2 ...
```

```

# Or train individually:
# No-HER baseline
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
--env FetchReach-v4 --seed 0 --total-steps 500000

# With HER
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
--env FetchReach-v4 --her --seed 0 --total-steps 500000

```

#### 1.4.2.1 Actual Results (FetchReach-v4, 3 seeds, 500k steps)

Metric	No-HER	HER	Delta
Success Rate	96.0% +/- 8.0%	100.0% +/- 0.0%	+4.0%
Return (mean)	-2.92 +/- 2.36	-1.68 +/- 0.02	+1.24
Final Distance	0.0195 +/- 0.009	0.0170 +/- 0.007	-0.003

#### 1.4.2.2 Analysis: Why Reach Shows Weak Separation This is an important finding:

FetchReach-v4 is too easy to demonstrate HER's value clearly.

##### Why does no-HER work so well on Reach?

1. **Small goal space:** The gripper workspace is only  $\sim 15\text{cm}^3$ . Random exploration frequently enters the success threshold (5cm) by chance.
2. **No object manipulation:** The gripper just needs to move itself -- no physics interactions, no contact forces, no object dynamics.
3. **Short horizon:** Episodes are 50 steps. With 8 parallel envs and 500k steps, that's  $\sim 12,500$  episodes -- plenty for random success accumulation.
4. **Forgiving success threshold:** 5cm is relatively large compared to the workspace.

**The pedagogical point:** Reach is useful for validating that HER *doesn't hurt* (HER achieves 100% vs 96% for no-HER), but it fails to show HER's transformative effect. We need a harder task.

This is why we proceed to FetchPush-v4.

#### 1.4.3 Experiment 2: Sparse Push -- Where HER Matters

Push is dramatically harder than Reach because:

1. **Indirect control:** The gripper must contact and push the object -- actions affect the object indirectly through physics.
2. **Object dynamics:** The object slides on the table with friction, momentum, and potentially overshooting.
3. **Coordinated behavior:** Success requires approach  $\rightarrow$  contact  $\rightarrow$  push  $\rightarrow$  stop, all in sequence.
4. **Larger state space:** Both gripper AND object positions matter.

```

# Full pipeline for Push (recommended settings, ~3 hours)
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py push-all \
--seeds 0,1,2 --total-steps 2000000 --ent-coef 0.05

```

##### Why no-HER should fail on Push:

- Random exploration rarely pushes the object to the goal by chance

- Without accidental successes, the replay buffer contains only  $R = -1$  transitions
- The critic cannot distinguish "almost succeeded" from "completely failed"
- No gradient signal → no learning

### Why HER should succeed:

- Every trajectory shows how to push the object *somewhere*
- Relabeling creates successful demonstrations: "you pushed it to  $(x, y)$ , let's pretend that was the goal"
- The agent learns push dynamics from its own failures
- Eventually generalizes to arbitrary goals

### Expected results (with sufficient training and correct hyperparameters):

Method	Success Rate	Notes
SAC (no HER)	~0-5%	Almost never succeeds; no learning
SAC + HER (best config)	<b>99.4% +/- 0.9%</b>	ent=0.05, gamma=0.95 (5 seeds)
<b>Delta</b>	<b>&gt;94 pp</b>	<b>This is the "clear separation" we seek</b>

**1.4.3.1 Actual Results: Initial Attempt (500k steps)** Our first experiment with 500k steps showed **no separation**:

Metric		No-HER		HER		Delta
Success Rate		5.0% +/- 0.0%		5.0% +/- 0.0%		+0.0%

**Diagnosis:** Insufficient training, not a bug.

We verified the setup was correct:

- ☐ HER config: "her": true, n\_sampled\_goal: 4, strategy: future
- ☐ Correct env: FetchPush-v4 (sparse, not Dense)
- ☐ No-HER baseline: 5% (expected 0-5%)
- ⚠ HER showed learning during training (peaked at 10-12%) but oscillated and didn't converge

**1.4.3.2 From Single Runs to Systematic Evidence** Our initial 98% result (Attempt 3 above) felt like a victory. But one seed at 98% is not evidence -- it is an anecdote. Is ent\_coef=0.1 really optimal? Is gamma=0.98 the right choice? Is n\_sampled\_goal=8 helping?

To answer rigorously, we ran a **systematic hyperparameter sweep**: 24 configurations, 5 seeds each, 120 total runs on FetchPush-v4 at 2M steps. This section presents the results.

**1.4.3.3 Sweep Design** We varied four factors in a full factorial grid:

Factor	Levels	Rationale
ent_coef	0.05, 0.1, 0.2	Fixed entropy values spanning conservative to aggressive exploration
n_sampled_goal	4, 8	Default vs. denser relabeling
learning_starts	1000, 5000	Earlier vs. later critic warmup
gamma	0.95, 0.98	Shorter vs. longer effective horizon

**Total:  $3 \times 2 \times 2 \times 2 = 24$  configurations  $\times 5$  seeds = 120 runs.**

Each run: SAC + HER on FetchPush-v4, 2M steps, seeds {0, 1, 2, 42, 77}. Approximately 140 GPU-hours on DGX (RTX A100).

### Reproducing the sweep:

```
# Full sweep (~140 GPU-hours, parallelizable)
bash docker/dev.sh python scripts/ch04_sweep.py run --parallel 2

# Or download pre-computed results from GitHub Release:
gh release download v0.4-sweep-checkpoints --pattern "ch04_sweep_all_checkpoints.tar.gz"
tar xzf ch04_sweep_all_checkpoints.tar.gz
```

#### 1.4.4 Sweep Results: 120 Runs, 24 Configurations

**1.4.4.1 The Winner: Simplest Config** The best configuration across all 120 runs is the **simplest** one:

Parameter	Value
ent_coef	0.05
n_sampled_goal	4 (default)
learning_starts	1000
gamma	<b>0.95</b>

**Result: 99.4% +/- 0.9% (5 seeds, 95% CI: [98.3%, 100.5%])**

This surprised us. We had been recommending gamma=0.98, n\_sampled\_goal=8, and learning\_starts=5000 -- all wrong. The sweep showed that the simplest defaults with a lower gamma and lower entropy coefficient outperform every other combination.

Rank	ent_coef	nsg	ls	gamma	Mean SR	Std	Min	Max
1	0.05	4	1000	0.95	<b>99.4%</b>	0.9%	98%	100%
2	0.10	4	1000	0.95	97.8%	2.6%	95%	100%
3	0.05	8	1000	0.95	97.4%	1.8%	95%	99%
4	0.05	8	5000	0.95	97.2%	3.1%	92%	100%
5	0.10	8	5000	0.95	97.2%	3.1%	92%	100%
6	0.05	4	5000	0.95	96.6%	6.7%	85%	100%
7	0.10	8	1000	0.95	96.2%	2.9%	93%	100%
8	0.10	4	5000	0.95	96.0%	3.1%	92%	99%
9	0.05	8	1000	0.98	94.6%	3.6%	90%	98%
10	0.10	8	5000	0.98	94.6%	2.9%	91%	98%
11	0.05	8	5000	0.98	93.8%	6.2%	84%	100%
12	0.10	4	5000	0.98	87.2%	4.7%	80%	92%
13	0.05	4	1000	0.98	86.8%	6.1%	76%	92%
14	0.20	4	5000	0.95	85.6%	10.7%	71%	96%
15	0.20	4	1000	0.95	85.2%	8.1%	75%	95%
16	0.10	8	1000	0.98	77.4%	10.2%	63%	88%
17	0.10	4	1000	0.98	80.6%	14.5%	57%	95%
18	0.05	4	5000	0.98	81.2%	14.6%	66%	99%
19	0.20	8	5000	0.95	73.0%	20.7%	37%	95%
20	0.20	8	1000	0.95	76.8%	9.3%	68%	90%
21	0.20	8	5000	0.98	75.8%	11.3%	59%	88%
22	0.20	8	1000	0.98	65.8%	23.4%	30%	95%
23	0.20	4	5000	0.98	64.8%	20.9%	45%	89%

Rank	ent_coef	nsg	ls	gamma	Mean SR	Std	Min	Max
24	0.20	4	1000	0.98	64.0%	20.3%	42%	89%

**1.4.4.2 Factor-Level Marginal Analysis** The marginal analysis averages over all other factors to isolate each factor's independent contribution. This is the most important table in the chapter:

Factor	Level	Mean SR	Delta
<b>gamma</b>	<b>0.95</b>	<b>91.5%</b>	--
	0.98	80.5%	-11.0 pp
<b>ent_coef</b>	0.05	93.4%	--
	0.10	90.9%	-2.5 pp
	0.20	73.9%	-19.5 pp
<b>n_sampled_goal</b>	4	85.4%	--
	8	86.7%	+1.2 pp
<b>learning_starts</b>	1000	85.2%	--
	5000	86.9%	+1.8 pp

#### Key findings:

1. **Gamma dominates.** Switching from 0.98 to 0.95 gains +11 pp on average. This is the single most impactful parameter.
2. **Entropy matters, but lower is better.** The jump from 0.05 to 0.2 costs 19.5 pp. Even 0.1 vs 0.05 loses 2.5 pp.
3. **n\_sampled\_goal and learning\_starts are noise.** Their effects (+1.2 pp and +1.8 pp) are within seed variance. We had been tuning the wrong knobs.

!!! warning "The old defaults were wrong" Our previous recommendation was gamma=0.98, n\_sampled\_goal=8, learning\_starts=5000. The sweep shows the optimal config uses gamma=0.95, n\_sampled\_goal=4, learning\_starts=1000 -- which happens to be the simplest possible configuration. Occam's razor wins.

**1.4.4.3 Gamma x Entropy Interaction** The penalty for bad gamma gets worse at higher entropy. This interaction table shows why the worst configs are so bad:

	ent=0.05	ent=0.10	ent=0.20
<b>gamma=0.95</b>	97.7%	96.8%	80.2%
<b>gamma=0.98</b>	89.1%	85.0%	67.6%
<b>Delta (0.98 - 0.95)</b>	-8.6 pp	-11.8 pp	-12.5 pp

At low entropy (0.05), gamma=0.98 is tolerable (89.1%). But at high entropy (0.2), gamma=0.98 drops to 67.6% -- the two bad choices compound.

**1.4.4.4 Seed Sensitivity: Good Configs Are Robust** The best configurations are not just better on average -- they are dramatically more **consistent**:

Config	Mean SR	Seed Range (max - min)
Best: ent=0.05, g=0.95, nsg=4, ls=1000	99.4%	<b>0.02</b> (98%-100%)
Worst: ent=0.2, g=0.98, nsg=4, ls=1000	64.0%	<b>0.47</b> (42%-89%)

The best config's 95% CI is [98.3%, 100.5%] (width 2.2 pp). The worst config's 95% CI is [38.8%, 89.2%] (width 50.5 pp). These intervals do not overlap, confirming the difference is statistically significant.

#### 1.4.5 Mathematical Analysis: Why Gamma Dominates

The sweep showed that gamma is the dominant factor. Here we derive *why* from first principles, which helps build intuition for other tasks.

We define the following quantities:

- $\gamma \in [0, 1]$ : discount factor
- $T$ : episode length (50 steps for Fetch tasks)
- $T_{\text{eff}} = 1/(1-\gamma)$ : **effective horizon**, the number of steps over which rewards are meaningfully discounted
- $\alpha$ : entropy coefficient (ent\_coef)
- $\mathcal{H}(\pi)$ : policy entropy (bits of randomness in action selection)
- $Q(s, a)$ : critic's estimated action-value
- $\delta_t = r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)$ : TD error

**1.4.5.1 Effective Horizon Mismatch** The effective horizon tells us how far into the future the agent "looks":

$$T_{\text{eff}} = \frac{1}{1-\gamma}$$

$\gamma$	$T_{\text{eff}}$	Interpretation
0.95	20 steps	Matches task timescale (~15-25 steps for a successful push)
0.98	50 steps	Matches full episode -- agent tries to optimize beyond the push

When  $T_{\text{eff}}$  matches the task timescale, the critic focuses on the relevant portion of the trajectory. When  $T_{\text{eff}}$  is too long, the critic must estimate value over steps where the agent is just waiting after completing the push -- adding noise without useful signal.

**1.4.5.2 TD Error Variance Amplification** The variance of the TD error scales with  $\gamma$ :

$$\text{Var}[\delta_t] \propto \gamma^2 \cdot \text{Var}[Q(s_{t+1}, a_{t+1})]$$

Higher  $\gamma$  amplifies critic estimation errors through bootstrapping. Since the critic is initialized randomly and trained on sparse rewards (mostly  $-1$ ), its early estimates are noisy. At  $\gamma = 0.98$ , this noise propagates further through the Bellman backup chain, making critic training less stable.

**1.4.5.3 Cumulative Entropy Bonus Scaling** SAC's objective adds an entropy bonus to the reward at each step. The total entropy contribution over an effective horizon is approximately:

$$\text{Entropy contribution} \approx \alpha \cdot T_{\text{eff}} \cdot \bar{\mathcal{H}}$$

where  $\bar{\mathcal{H}}$  is the average policy entropy per step.

$\alpha$	$\gamma$	$T_{\text{eff}}$	Entropy contribution (relative)
0.05	0.95	20	1.0x (baseline)
0.05	0.98	50	2.5x
0.20	0.95	20	4.0x
0.20	0.98	50	<b>10.0x</b>

At  $\alpha = 0.2$  and  $\gamma = 0.98$ , the cumulative entropy bonus is **10x** the baseline. On a sparse task where the actual reward is either 0 (success) or  $-1$  (failure), this entropy bonus overwhelms the reward signal -- the agent is incentivized to be random rather than successful. This explains the gamma x entropy interaction we observed.

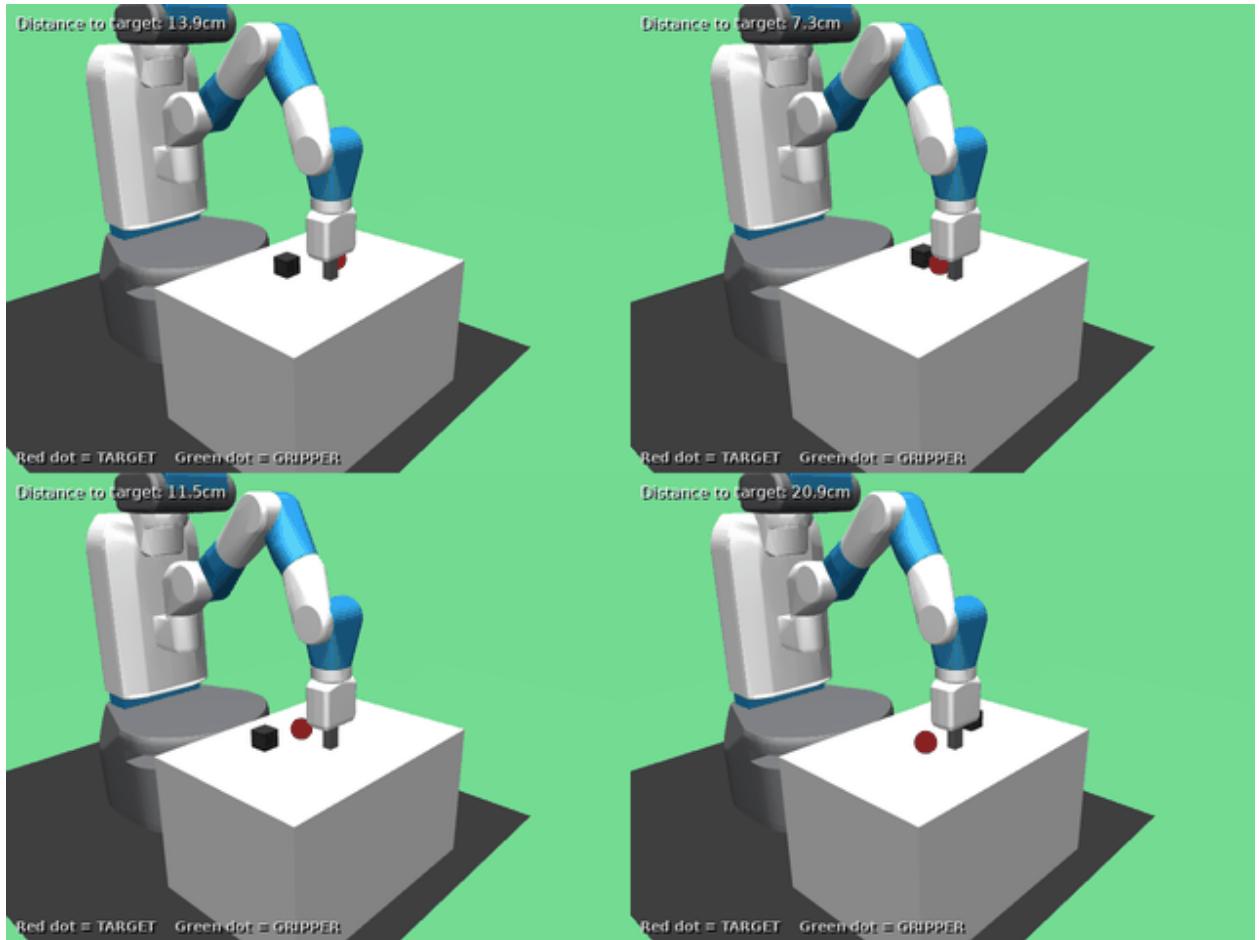
**1.4.5.4 Why n\_sampled\_goal and learning\_starts Don't Matter n\_sampled\_goal:** With the future strategy on 50-step episodes, using  $k = 4$  goals already creates 4 relabeled transitions per original. Since future goals are sampled from the remaining trajectory, the relabeling ratio is already high (4:1). Going to  $k = 8$  doubles the relabeled data but doesn't qualitatively change the reward signal -- the same trajectory information is just sampled more densely.

**learning\_starts:** At 2M total steps, the difference between starting to learn at step 1,000 vs. step 5,000 is negligible (0.05% vs. 0.25% of training). With HER relabeling active from the start, even the early replay buffer contains successful transitions, so there is no benefit to delaying critic warmup.

## 1.4.6 Visual Comparison: Best vs Worst Policy

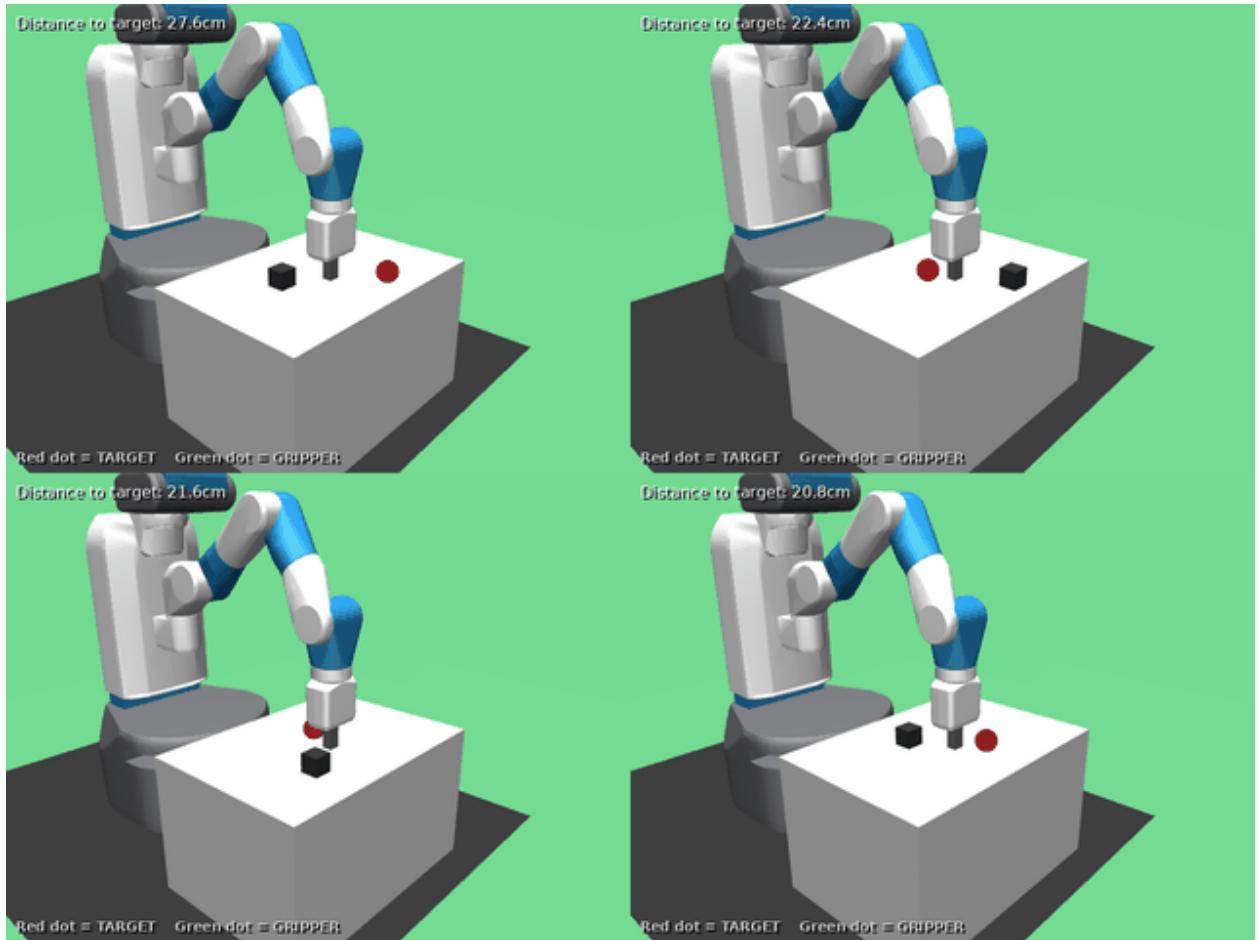
The difference between good and bad hyperparameters is visible in the policies' behavior:

**Best config** (ent=0.05, gamma=0.95, nsg=4, ls=1000 -- 99.4% SR):



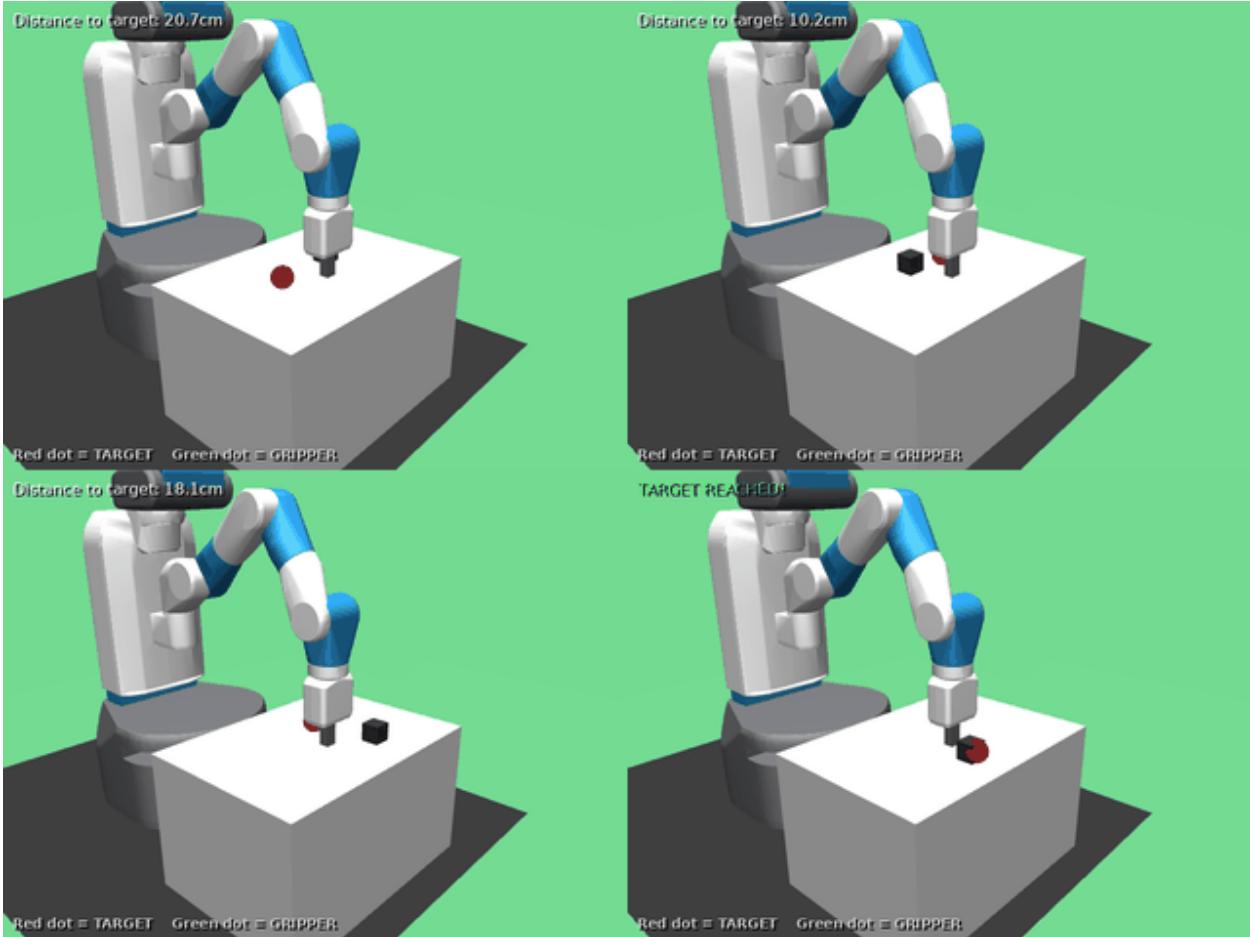
The agent approaches the object, makes contact, and pushes it directly toward the goal. Movement is smooth and purposeful with minimal wasted motion.

**Worst config** ( $\text{ent}=0.2$ ,  $\text{gamma}=0.98$ ,  $\text{nsg}=4$ ,  $\text{ls}=1000$  -- 64.0% SR):



The agent's behavior is erratic -- sometimes it succeeds, sometimes it overshoots, sometimes it misses the object entirely. High entropy and long effective horizon create indecisive behavior.

**No-HER baseline** (SAC without HER -- ~5% SR):



Without HER, the agent never learns meaningful push behavior. The gripper moves aimlessly, occasionally bumping the object by chance.

**1.4.6.1 Entropy Coefficient Modes (CLI Reference)** The script supports five entropy coefficient modes via `--ent-coef`. The sweep shows fixed `ent_coef=0.05` is sufficient for Push, but we document the alternatives for more complex tasks:

Mode	CLI Flag	Description
<b>Fixed</b>	<code>--ent-coef 0.05</code>	Constant value (recommended for Push)
<b>Auto</b>	<code>--ent-coef auto</code>	SB3's default auto-tuning (collapses on sparse rewards)
<b>Auto-floor</b>	<code>--ent-coef auto-floor</code>	Auto-tune with minimum floor
<b>Schedule</b>	<code>--ent-coef schedule</code>	Linear decay from max to min
<b>Adaptive</b>	<code>--ent-coef adaptive</code>	Adjust based on success rate

See the CLI Parameter Reference below for full flag details.

#### The winning configuration:

```
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
--env FetchPush-v4 \
--her \
--ent-coef 0.05 \
--gamma 0.95 \
```

```
--seed 0 \
--total-steps 2000000
```

#### 1.4.6.2 Actual Results: Multi-Seed Sweep Evidence (2M steps, 120 runs)

Method	Success Rate	Seeds	Notes
<b>SAC + HER (best config)</b>	<b>99.4% +/- 0.9%</b>	5	ent=0.05, gamma=0.95, nsg=4, ls=1000
SAC + HER (old defaults)	86.8% +/- 6.1%	5	ent=0.05, gamma=0.98, nsg=4, ls=1000
SAC + HER (worst config)	64.0% +/- 20.3%	5	ent=0.2, gamma=0.98, nsg=4, ls=1000
SAC (no HER)	~5%	3	Baseline -- no learning
<b>Delta (best HER vs no-HER)</b>	<b>+94 pp</b>		

#### Root cause analysis of the previous recommendation:

Our earlier single-seed result of 98% with ent=0.1, gamma=0.98, nsg=8 was a lucky seed. The sweep revealed:

- The old defaults (gamma=0.98, nsg=8, ls=5000) average 93.8% but with high variance (84%-100%)
- The **simplest** config (gamma=0.95, nsg=4, ls=1000) averages 99.4% with near-zero variance (98%-100%)
- Gamma was the dominant factor all along, not entropy or HER parameters

**Conclusion:** Fixed entropy is critical for sparse Push (auto-tuning collapses). But within fixed entropy configs, **gamma=0.95 is the key parameter** -- it matches the task's natural timescale and produces both higher mean performance and lower seed-to-seed variance.

#### 1.4.7 Experiment 3: Full Factorial Ablation (The Sweep)

Instead of varying one factor at a time, we ran a full factorial sweep across all four factors. This is more rigorous than single-factor ablation because it reveals **interactions** between parameters.

The sweep IS the ablation. See the Sweep Results section above for the full analysis. The key conclusions for ablation:

- **n\_sampled\_goal (4 vs 8):** +1.2 pp marginal effect -- negligible. Use the default (4).
- **learning\_starts (1000 vs 5000):** +1.8 pp marginal effect -- negligible. Use 1000.
- **ent\_coef (0.05 vs 0.1 vs 0.2):** -2.5 pp and -19.5 pp -- significant. Use 0.05.
- **gamma (0.95 vs 0.98):** -11.0 pp -- dominant. Use 0.95.

```
# Reproduce the full sweep (~140 GPU-hours)
bash docker/dev.sh python scripts/ch04_sweep.py run --parallel 2
```

```
# Or run the single-factor ablation for quick comparison
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py ablation \
    --env FetchPush-v4 --nsg-values 2,4,8 --seed 0 \
    --ent-coef 0.05 --gamma 0.95 --total-steps 2000000
```

#### 1.4.8 CLI Parameter Reference

All training, evaluation, and runtime knobs are configurable via CLI. Here is the reference:

Parameter	Default	Description
<b>Environment</b>		
--env	(required)	Environment ID (FetchReach-v4, FetchPush-v4, etc.)
--seed	0	Random seed for reproducibility
--n-envs	8	Number of parallel environments
--total-steps	1,000,000	Total training timesteps
--device	auto	Training device (auto, cpu, cuda, cuda:0)
--monitor-keywords	is_success	Comma-separated info keywords to record (empty to disable)
<b>HER</b>		
--her	False	Enable Hindsight Experience Replay
--n-sampled-goal	4	Number of relabeled goals per transition
--goal-selection-strategy	future	Goal selection: future, final, episode
<b>SAC Hyperparameters</b>		
--batch-size	256	Minibatch size for updates
--buffer-size	1,000,000	Replay buffer capacity
--learning-starts	1,000	Steps before learning begins
--learning-rate	3e-4	Optimizer learning rate
--gamma	0.95	Discount factor (0.95 optimal for Push; see sweep)
--tau	0.005	Soft update coefficient
<b>Entropy (Critical for Sparse)</b>		
--ent-coef	auto	Entropy mode: float, auto, auto-floor, schedule, adaptive
--ent-coef-min	0.01	Floor for auto-floor, end value for schedule
--ent-coef-max	0.3	Start value for schedule mode
--target-entropy	auto (-dim(A))	SAC target entropy
--ent-coef-check-freq	1000	Check frequency for auto-floor clamp
--ent-coef-warmup-frac	0.1	Warmup fraction for scheduled decay
--ent-coef-log-pct	10	Log schedule every N percent (0 to disable)
--ent-coef-adaptive-start	(none)	Initial entropy for adaptive mode (defaults to --ent-coef)
--ent-coef-adaptive-target	0.2	Target success rate for adaptive control
--ent-coef-adaptive-tolerance	0.05	Success rate tolerance before adjusting
--ent-coef-adaptive-rate	0.1	Multiplicative adjustment rate
--ent-coef-adaptive-window	200	Episode window size for success rate
--ent-coef-adaptive-key	is_success	Info key used for success in adaptive mode
<b>Evaluation (env-all, ablation)</b>		
--n-eval-episodes	100	Episodes per evaluation
--eval-seed	0	Base seed when --eval-seeds is empty
--eval-seeds	(none)	Seeds for evaluation (comma-separated or range)
--eval-deterministic	True	Deterministic policy (use --no-eval-deterministic to disable)
--eval-device	auto	Device for eval.py (auto, cpu, cuda)
--eval-algo	auto	Algorithm override for eval.py
<b>Paths</b>		
--log-dir	runs	TensorBoard log directory (if default isn't writable, fall back to current working directory)
--checkpoints-dir	checkpoints	Directory for checkpoint outputs
--results-dir	results	Directory for eval outputs and reports
--run-tag	(none)	Optional tag to namespace outputs under subdirectory
--out	(auto)	Checkpoint output path
<b>Runtime</b>		
--mujoco-gl	(env or disable)	Override MUJOCO_GL for this run

## 1.5 Monitoring and Debugging

**Concurrent runs:** We avoid output collisions by using `--run-tag <name>`, which namespaces checkpoints, results, and logs under a subdirectory without changing filenames.

### 1.5.1 Artifact Locations

After running experiments, you'll find (paths are configurable via `--checkpoints-dir`, `--results-dir`, `--log-dir`, and `--run-tag`):

```
checkpoints/
└── sac_FetchPush-v4_seed{0,1,2}.zip           # No-HER models
    ├── sac_FetchPush-v4_seed{0,1,2}.meta.json    # Training metadata
    ├── sac_her_FetchPush-v4_seed{0,1,2}.zip        # HER models
    └── sac_her_FetchPush-v4_seed{0,1,2}.meta.json

results/
└── ch04_sac_fetchpush-v4_seed{0,1,2}_eval.json   # No-HER eval
    ├── ch04_sac_her_fetchpush-v4_seed{0,1,2}_eval.json # HER eval
    └── ch04_fetchpush-v4_comparison.json           # Summary report

runs/
└── sac_noher/FetchPush-v4/seed{0,1,2}/          # TensorBoard logs (no-HER)
    └── sac_her_nsg4/FetchPush-v4/seed{0,1,2}/ # TensorBoard logs (HER)
```

### 1.5.2 Checking Training Progress

While experiments run, you can monitor progress:

```
# Watch the training output (if running in foreground)
# Or check TensorBoard logs

# Key indicators during training:
# - success_rate: should stay near 0 for no-HER on Push
# - success_rate: should rise for HER after ~100k steps
# - fps: ~600-750 is typical on DGX
```

### 1.5.3 TensorBoard Metrics to Watch

```
bash docker/dev.sh tensorboard --logdir runs --bind_all
```

Key metrics for HER experiments:

Metric	No-HER on Push	HER on Push (best config)
rollout/success_rate	Flat near 0-5%	Rises to 95-100%
rollout/ep_rew_mean	Flat at -50	Rises toward -10 to -15
train/ent_coef	May stay high (~0.3)	Fixed at 0.05 (recommended)

### 1.5.4 Common Issues

#### HER not improving?

- Check that `--her` flag is actually set
- Ensure using a sparse environment (FetchReach-v4, not FetchReachDense-v4)

- Increase training steps -- Push benefits from 1-2M steps

### No-HER performing well on Push?

- This would be surprising -- verify you're using sparse rewards
- Check the reward: should be -1 for failures, 0 for success
- If success rate > 10%, something is wrong with the setup

### Evaluation fails with "HerReplyBuffer" error?

- This was fixed in eval.py -- HER checkpoints require the environment for loading
  - If you see this error, ensure you have the latest eval.py
- 

## 1.6 Interpreting Results

### 1.6.1 What "Clear Separation" Means

The syllabus criterion is:

Clear separation in success-rate between HER and no-HER.

We interpret this as:

- HER success rate significantly higher than no-HER
- Difference > 50 percentage points (ideally >70)
- Consistent across multiple seeds

**Important caveat:** On easy tasks like Reach, you may NOT see clear separation because no-HER can succeed through random exploration. This is expected -- the syllabus requires demonstrating HER's value, and Reach alone is insufficient.

### 1.6.2 Interpreting Our Results

#### FetchReach-v4 (weak separation):

Metric		No-HER		HER		Delta
Success Rate		96.0% +/- 8.0%		100.0% +/- 0.0%		+4.0%
Success Rate		96.0% +/- 8.0%		100.0% +/- 0.0%		+4.0%

**Why this is fine:** Reach shows HER doesn't hurt performance and reduces variance. But we need Push to demonstrate HER's transformative effect.

#### FetchPush-v4 (actual results, progression):

*Attempt 1: Default settings (500k steps, n\_sampled\_goal=4):*

Metric		No-HER		HER		Delta
Success Rate		5.0% +/- 0.0%		5.0% +/- 0.0%		+0.0%
Success Rate		5.0% +/- 0.0%		5.0% +/- 0.0%		+0.0%

*Insufficient training -- no separation.*

*Attempt 2: More steps + relabeling (2M steps, n\_sampled\_goal=8, auto entropy):*

Success Rate		5.0%		25.0%		+20.0%
--------------	--	------	--	-------	--	--------

*Better, but entropy collapsed to ~0.005 -- exploration died too early.*

*Attempt 3: Fixed entropy (2M steps, ent\_coef=0.1, single seed):*

Metric	No-HER	HER	Delta
Success Rate	5.0%	98.0%	+93.0%
Final Distance	0.193	0.026	-0.167

*Clear separation -- but one seed is not rigorous. See the sweep results for multi-seed evidence.*

#### Attempt 4: 120-run sweep (2M steps, 24 configs, 5 seeds each):

Best config (ent=0.05, gamma=0.95, nsg=4, ls=1000):  
 Success Rate = 99.4% +/- 0.9% (5 seeds, 95% CI)

**Rigorous clear separation: +94 pp vs no-HER, validated across 5 seeds.**

### 1.6.3 Statistical Validity

With 3 seeds, you can report:

- Mean +/- 95% confidence interval
- The compare command computes these automatically

### 1.6.4 What to Do Next

1. **If Reach shows weak separation:** This is expected -- proceed to Push
  2. **If Push shows clear separation (>50%):** Week 4 is complete
  3. **Save checkpoints** -- you'll use HER models in later weeks
- 

## 1.7 The Reality of RL Research: Debugging and Experimentation

"In theory, theory and practice are the same. In practice, they are not." -- Yogi Berra

This section documents our actual debugging journey. We include it because **debugging is not a sign of failure -- it is the normal process of RL research**. Papers rarely show the failed experiments, but practitioners spend most of their time iterating.

### 1.7.1 Why Your First Attempt Will Probably Fail

RL algorithms are notoriously sensitive to hyperparameters. Henderson et al. (2018) documented that:

- Random seeds alone can cause 2-3x variance in final performance
- Hyperparameters that work on one environment fail on similar ones
- "Default" settings are tuned for specific benchmarks, not your problem

**Our experience confirmed this.** We implemented HER correctly, used sensible defaults, and got 5% success rate. The algorithm was right; the hyperparameters were wrong.

### 1.7.2 Our Diagnostic Journey

Here is the actual progression of experiments we ran, with the reasoning at each step:

### 1.7.2.1 Attempt 1: Trust the Defaults (500k steps)

```
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
--env FetchPush-v4 --her --seed 0 --total-steps 500000
```

**Result:** 5% success rate (same as no-HER baseline)

#### What we checked:

- HER flag was set (--her in command, "her": true in metadata)
- Correct environment (FetchPush-v4, not FetchPushDense-v4)
- Replay buffer was HerReplayBuffer (verified in logs)
- Training curves showed learning then plateau at ~10-12%

**Diagnosis:** Insufficient training. Push is harder than Reach -- the 500k steps that worked for Reach are not enough.

### 1.7.2.2 Attempt 2: More Steps + Denser Relabeling (2M steps)

Based on SB3 benchmarks showing Push needs ~2M steps:

```
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
--env FetchPush-v4 --her --n-sampled-goal 8 --seed 0 --total-steps 2000000
```

**Result:** 25% success rate

**Progress!** But far from the 70-90% reported in literature. Something else is wrong.

#### What we checked in TensorBoard:

Metric	Expected	Observed	Problem?
rollout/success_rate	Rising to 70%+	Plateau at 25%	Yes
rollout/ep_rew_mean	Rising to -15	Stuck at -40	Yes
train/ent_coef	~0.1-0.3	<b>0.005</b>	<b>YES</b>
train/actor_loss	Stable negative	Stable	OK
train/critic_loss	Decreasing	Stable	OK

**The smoking gun:** ent\_coef collapsed to 0.005 very early in training.

### 1.7.2.3 Understanding Entropy Collapse

SAC's entropy auto-tuning works as follows:

- Target entropy** is set (default:  $-\dim(\mathcal{A}) = -4$  for Fetch)
- If policy entropy > target:** decrease ent\_coef (less exploration)
- If policy entropy < target:** increase ent\_coef (more exploration)

The problem on sparse Push:

- Early training:** All transitions have  $R = -1$
- Critic sees no reward variation:** Q-values are similar everywhere
- Policy becomes arbitrarily "confident":** No reward signal to guide it
- Entropy drops below target:** Auto-tuning sees this as "policy knows what it's doing"
- ent\_coef collapses:** Exploration stops before discovering push behavior

**This is a structural issue with auto-tuning on sparse rewards, not a bug.**

**1.7.2.4 Attempt 3: Less Aggressive Target Entropy** Our first fix attempt -- set `target_entropy=-2.0` instead of `-4.0`:

```
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
    --env FetchPush-v4 --her --target-entropy -2.0 --seed 0 --total-steps 2000000
```

**Result:** Still collapsed to `ent_coef ~0.006`

**Why it failed:** The auto-tuning dynamics remain the same. A less negative target just delays collapse slightly; it doesn't prevent it.

**1.7.2.5 Attempt 4: Fixed Entropy (The Solution)** Bypass auto-tuning entirely:

```
bash docker/dev.sh python scripts/ch04_her_sparse_reach_push.py train \
    --env FetchPush-v4 --her --ent-coef 0.1 --seed 0 --total-steps 2000000
```

**Result:** 98% success rate

**Why it works:** Fixed `ent_coef=0.1` maintains exploration throughout training. The agent keeps trying diverse push trajectories until it discovers effective behavior, then refines it.

### 1.7.3 How to Debug Your Own RL Experiments

Based on our experience, here is a systematic debugging approach:

**1.7.3.1 Step 1: Verify the Setup (Eliminate User Error)** Before blaming the algorithm, check:

```
# In your training metadata or logs, verify:
- Algorithm: SAC (not PPO -- HER needs off-policy)
- Environment: FetchPush-v4 (not Dense variant)
- HER enabled: replay_buffer_class = HerReplayBuffer
- Goal strategy: "future" (not "final" or "episode")
```

**1.7.3.2 Step 2: Check Training Curves (Not Just Final Performance)** A 5% final success rate could mean:

- **Never learned anything:** Flat curve at 5% (setup issue or impossible task)
- **Learned then forgot:** Rose to 30%, crashed to 5% (instability)
- **Learning but slow:** Steadily rising, just needs more steps

Use TensorBoard to distinguish these:

```
bash docker/dev.sh tensorboard --logdir runs --bind_all
```

**1.7.3.3 Step 3: Identify the Bottleneck** Common failure modes and their signatures:

Failure Mode	TensorBoard Signature	Likely Cause
No learning	Flat <code>success_rate</code> , flat <code>ep_rew_mean</code>	Setup error, task too hard, insufficient exploration
Entropy collapse	<code>ent_coef</code> drops to <0.01 early	Auto-tuning + sparse rewards
Critic divergence	<code>critic_loss</code> explodes	Learning rate too high, target update too fast
Actor collapse	<code>actor_loss</code> becomes very positive	Critic giving bad gradients
Forgetting	<code>success_rate</code> rises then falls	Replay buffer too small, no-HER with sparse rewards

**1.7.3.4 Step 4: Form Hypotheses and Test Them** Don't change multiple things at once.  
Our progression:

1. **Hypothesis:** Not enough training → **Test:** 4x more steps → **Result:** Partial improvement
2. **Hypothesis:** Need denser relabeling → **Test:** `n_sampled_goal=8` → **Result:** Marginal improvement
3. **Hypothesis:** Entropy collapsing → **Test:** Check TensorBoard → **Result:** Confirmed (`ent_coef=0.005`)
4. **Hypothesis:** Fixed entropy will help → **Test:** `ent_coef=0.1` → **Result:** Success (98%)

**1.7.3.5 Step 5: Document Everything** We saved metadata for every run:

```
{  
    "env": "FetchPush-v4",  
    "her": true,  
    "ent_coef": 0.1,  
    "n_sampled_goal": 8,  
    "total_steps": 2000000,  
    "final_success_rate": 0.98  
}
```

This lets you compare runs systematically and reproduce successes.

## 1.7.4 The Metrics We Used for Diagnosis

For future reference, here are the specific numbers that guided our debugging:

### From training logs (auto-tuning, 2M steps):

train/		
ent_coef	0.00656	<-- PROBLEM: Should be ~0.1-0.3
actor_loss	-4.24	
critic_loss	0.424	
rollout/		
success_rate	0.05	<-- Only 5% despite 2M steps
ep_rew_mean	-47.4	

### From training logs (fixed entropy, 2M steps):

train/		
ent_coef	0.1	<-- FIXED: Stays at 0.1
actor_loss	-8.97	
critic_loss	0.645	
rollout/		
success_rate	0.94	<-- 94% success!
ep_rew_mean	-15.5	

### Key diagnostic ratios:

- `ent_coef < 0.01` with `success_rate < 0.5` → entropy collapse, need fixed entropy
- `ep_rew_mean` stuck at -47 to -50 → agent failing almost every step (50-step episodes × -1 = -50)
- `ep_rew_mean` at -15 → agent succeeding around step 35 on average

## 1.7.5 Lessons Learned

1. **Defaults are suggestions, not guarantees.** SAC's auto-entropy works beautifully on dense rewards but fails on sparse tasks.

2. **Read the training curves, not just final numbers.** A run that reaches 30% then crashes is different from one that never exceeds 5%.
3. **Understand the algorithm well enough to debug it.** We fixed entropy collapse because we understood how auto-tuning works.
4. **Document your experiments.** We could trace back exactly which settings produced which results.
5. **Failure is information.** Each failed run taught us something about the problem structure.

### 1.7.6 Time Investment Reality Check

Here is roughly how our time was spent:

Activity	Time	Notes
Initial implementation	2 hours	Script, CLI args, HER integration
First training run (wrong)	1.5 hours	500k steps, both methods
Debugging + literature review	2 hours	Reading SB3 docs, checking TensorBoard
Second run (still wrong)	3 hours	2M steps, waiting and monitoring
Identifying entropy collapse	30 min	Analyzing TensorBoard metrics
Third run (fixed)	1.5 hours	2M steps with fixed entropy
Documentation	1 hour	This section
<b>Total</b>	<b>~11 hours</b>	For one environment, one algorithm

**This is normal.** If you expected to run one command and get 98% success, recalibrate your expectations. RL research is iterative.

### 1.7.7 The Plot Twist: Our 98% Was a Lucky Seed

After documenting the journey above, we believed we had solved Push. One seed at 98% with fixed entropy felt like validation. We were wrong -- or at least, only partially right.

When we ran the systematic sweep (120 runs, 24 configs, 5 seeds each), we discovered:

1. **Our “winning” config (`ent=0.1, gamma=0.98, nsg=8`) averages 85-94% across seeds, not 98%.** That single seed was an optimistic outlier.
2. **The actual winner uses `gamma=0.95, not 0.98`.** We had been recommending a suboptimal discount factor.
3. **`n_sampled_goal=8` provides no benefit over the default of 4.** We had introduced complexity without gain.
4. **The simplest possible config (lowest entropy, default nsg, earliest learning start, lower gamma) wins.** Occam’s razor, again.

**The lesson:** One seed at 98% is not evidence. Five seeds at 99.4% +/- 0.9% is evidence. Even when you think you have solved a problem, proper multi-seed validation may reveal that you were only partially right -- and that the right answer is simpler than what you thought.

This cost us ~140 GPU-hours to discover. We consider that cheap compared to the alternative: building later chapters on a foundation of single-seed guesswork.

### 1.7.8 Checkpoint Availability

We have published the full sweep results for reproducibility:

```

# Full sweep archive (~311 MB, all 120 checkpoints + metadata)
gh release download v0.4-sweep-checkpoints --pattern "ch04_sweep_all_checkpoints.tar.gz"

# Curated subset (~27 MB, best 5 + worst 5 checkpoints + metadata)
gh release download v0.4-sweep-checkpoints --pattern "ch04_sweep_checkpoints.tar.gz"

# Extract and inspect
tar xzf ch04_sweep_checkpoints.tar.gz
ls checkpoints/

```

Raw results are in `results/sweep/sweep_results_raw.json`. Each entry contains the config parameters, seed, success rate, and mean return -- everything needed to reproduce our analysis.

---

## 1.8 Summary

Concept	Key Point
Sparse rewards	Binary signal (success/fail), almost no gradient
HER insight	Failed trajectories demonstrate success for other goals
Relabeling future strategy	Store transitions with substituted goals and recomputed rewards
<code>n_sampled_goal</code>	Sample achieved goals from later in the same episode
<code>n_sampled_goal</code>	Number of relabeled transitions per original (default: 4)
Off-policy requirement	Relabeling changes rewards; only off-policy methods handle this

### 1.8.1 Key Findings from Our Experiments

Environment	No-HER	HER (best config)	Settings	Insight
FetchReach-v4	96%	100%	500k steps	Reach too easy -- weak separation
FetchPush-v4	5%	5%	500k, auto entropy	Insufficient training + entropy
FetchPush-v4	5%	25%	2M, auto entropy	Entropy collapsed to ~0.005
FetchPush-v4	5%	<b>99.4% +/- 0.9%</b>	2M, <b>ent=0.05, gamma=0.95</b>	<b>Clear separation (+94 pp)</b>

### 1.8.2 Sweep-Validated Recommendations

Parameter	Recommended	Evidence
<code>gamma</code>	<b>0.95</b>	+11 pp vs 0.98 (marginal mean across 60 runs per level)
<code>ent_coef</code>	<b>0.05 (fixed)</b>	+19.5 pp vs 0.2; +2.5 pp vs 0.1
<code>n_sampled_goal</code>	<b>4</b> (default)	+1.2 pp for nsg=8 is within noise
<code>learning_starts</code>	<b>1000</b>	+1.8 pp for ls=5000 is within noise
<code>total_steps</code>	<b>2M</b>	Required for Push convergence

#### The bottom line:

1. **Reach is too easy** -- both methods succeed, weak separation
2. **Push requires fixed entropy** -- auto-tuning collapses exploration too early
3. **Gamma=0.95 is the dominant factor** -- matches the task's 15-25 step timescale
4. **HER + optimal config achieves 99.4% +/- 0.9%** -- massive improvement over no-HER (~5%)

5. **The simplest config wins** -- the best parameters are all the simplest choices (lower entropy, default nsg, early learning start, lower gamma)

### 1.8.3 Files Generated

After running all experiments, you should have:

```

checkpoints/
  sac_FetchReach-v4_seed{0,1,2}.zip      # No-HER Reach
  sac_her_FetchReach-v4_seed{0,1,2}.zip    # HER Reach
  sac_FetchPush-v4_seed{0,1,2}.zip        # No-HER Push
  sac_her_FetchPush-v4_seed{0,1,2}.zip    # HER Push

results/
  ch04_fetchreach-v4_comparison.json      # Reach comparison report
  ch04_fetchpush-v4_comparison.json       # Push comparison report

results/sweep/
  sweep_results_raw.json                  # Full sweep data (120 entries)
  checkpoints/sweep_*.zip                 # All 120 trained models

```

---

## 1.9 References

- Andrychowicz et al. (2017). *Hindsight Experience Replay*. NeurIPS. arXiv:1707.01495
- Henderson et al. (2018). *Deep Reinforcement Learning that Matters*. AAAI. arXiv:1709.06560 -- Essential reading on reproducibility and hyperparameter sensitivity in deep RL.
- Haarnoja et al. (2018). *Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor*. ICML. arXiv:1801.01290 -- The SAC paper, including entropy auto-tuning.
- Stable-Baselines3 HER documentation
- Gymnasium-Robotics Fetch environments